

# Techtalk

Gluster in 7d

October 07, 2016

## What's Gluster?



Figure 1:

“GlusterFS is a scalable network filesystem. Using common off-the-shelf hardware,

you can create large, distributed storage solutions for media streaming, data analysis, and other data- and bandwidth-intensive tasks.”

– <https://www.gluster.org/>

## Where do we use it?

- Mediapool 777 & 888
- Supplier uploads
- Ingestion Share (new starfox)

## Why do we use it?

- Redundancy on the cheap in the pre GCS world
- Less complex than legacy mediapools\*
- *\*In Theory*

## Simple use case

- Supplier uploads

supplier-upload.7digital.net.ssh																						
	Queue			Session rate			Sessions					Bytes		Denied		Errors		Warnings		Status	LastChk	
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp			Retr
Frontend				0	8	-	1	34	2 000	18 061			5 547 475 690	52 116 694	0	0	0				OPEN	
ingestion04.nix.sys.7d	0	0	-	0	4	0	17	-	8 926	8 926	5s		1 904 065 340	22 746 569	0		0	10	0	2h8m UP	L4OK in 0ms	
ingestion05.nix.sys.7d	0	0	-	0	4	1	18	-	9 096	9 096	2s		3 643 410 350	29 370 125	0		0	11	0	27d2h UP	L4OK in 1ms	
Backend	0	0		0	8	1	34	200	18 061	18 022	2s		5 547 475 690	52 116 694	0	0	0	21	0	27d2h UP		

- root@ingestion05:~# gluster vol info  
Volume Name: supplier-upload  
Type: Replicate  
Volume ID: 01cec521-7331-4b29-a90e-e23a823297e3  
Status: Started  
Number of Bricks: 1 x 2 = 2  
Transport-type: tcp  
Bricks:  
Brick1: ingestion04:/export/brick00/supplier-upload  
Brick2: ingestion05:/export/brick00/supplier-upload

## What it looks like

- ops@ingestion05:~\$ tail -2 /proc/mounts  
/dev/mapper/debian-brick00 /export/brick00 ext4 \

```
rw,noatime,user_xattr,barrier=1,data=ordered 0 0
ingestion05:/supplier-upload /srv fuse.glusterfs \
rw,relatime,user_id=0,group_id=0,default_permissions,allow_other,max_read=131072 0 0
```

- [ftpadmin]  
comment = FTP Share Directory  
path = /srv/ftp

## What it looks like (cont.)

- root@ingestion05:/export/brick00/supplier-upload/ftp# ls -t | head  
tunecore  
theorchard  
consolidated  
FUGA  
fuga\_guvera  
kontor  
sonyset  
believe  
sonyhd  
labelworx  
root@ingestion05:/srv/ftp# ls -t | head  
theorchard  
kontor  
tunecore  
consolidated  
FUGA  
fuga\_guvera  
sonyset  
believe  
sonyhd  
labelworx

## More complex use case

- mp777

hms-gluster01.prod.hms01







































	Options			Session rate			Sessions			In	Bytes			Dequeued			Errors			Warnings			Status	LastChk
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit		Cur	Max	Limit	Req	Resp	Req	Conn	Resp	Req	Conn	Resp	Req		
Frontend	0	0	-	1	140	-	1	140	20700	0	7	228	784	0	0	0	0	0	0	0	0	0	OK	
prod-hms01	0	0	-	1	73	-	1	110	-	3 381 895	3 381 575	14	613 321 905	0	0	0	102	357	320	0	27628	UP	1703625501.111816	
prod-hms02	0	0	-	1	73	-	1	77	-	3 381 870	3 381 560	14	603 462 490	0	0	0	102	357	320	0	27628	UP	1703625501.111816	
Backend	0	0	-	1	140	-	1	140	1 000	0	7	228	784	0	0	0	102	357	320	0	27628	UP	1703625501.111816	

- ops@prod-hms01:~\$ df -h  
Filesystem                      Size    Used Avail Use% Mounted on  
rootfs                           7.3G   1.5G   5.5G   21% /

udev	10M	0	10M	0%	/dev
tmpfs	397M	2.2M	395M	1%	/run
/dev/mapper/debian-root	7.3G	1.5G	5.5G	21%	/
tmpfs	5.0M	0	5.0M	0%	/run/lock
tmpfs	794M	0	794M	0%	/run/shm
/dev/vda1	228M	20M	197M	10%	/boot
ctr-prod-hms-b00:/gv01	610T	606T	3.4T	100%	/srv/gv01

## Complex (cont.)

- ops@prod-hms01:~\$ grep gv01 /etc/fstab  
ctr-prod-hms-b00:/gv01 /srv/gv01 glusterfs \  
defaults,\_netdev,backupvolfile-server=ctr-prod-hms-b01 0 0

<u>Name</u>	<u>Type</u>	<u>Priority</u>	<u>Content</u>
 ctr-prod-hms-B00.nix.sys.7d	A		10.112.16.10
 ctr-prod-hms-b01.nix.sys.7d	A		10.112.16.11
 ctr-prod-hms-d00.nix.sys.7d	A		10.112.16.12
 ctr-prod-hms-d01.nix.sys.7d	A		10.112.16.13
 gs2-prod-hms-A00.nix.sys.7d	A		10.108.16.10
 gs2-prod-hms-a01.nix.sys.7d	A		10.108.33.29
 gs2-prod-hms-a02.nix.sys.7d	A		10.108.16.18
 gs2-prod-hms-c00.nix.sys.7d	A		10.108.16.19
 gs2-prod-hms-c01.nix.sys.7d	A		10.108.16.20
 hms-hub00.nix.sys.7d	A		10.120.3.28
 mediapool777-wr.nix.sys.7d	CNAME		prod-hms-wr00.nix.sys.7d
 mediapool888-wr.nix.sys.7d	CNAME		prod-hms-wr01.nix.sys.7d
 mediapool999-wr.nix.sys.7d	CNAME		ofc-prod-hms-wr00.nix.sys.7d
 metadata-store.prod.svc.7d	CNAME		prod-hms-metadata-storage01.nix.sys.7d
 ofc-hms88.nix.sys.7d	A		10.120.19.15
 ofc-prod-hms-a00.nix.sys.7d	A		10.120.23.11
 ofc-prod-hms-a01.nix.sys.7d	A		10.120.23.14
 ofc-prod-hms-a02.nix.sys.7d	A		10.120.23.17
 ofc-prod-hms-b00.nix.sys.7d	A		10.120.23.10
 ofc-prod-hms-b01.nix.sys.7d	A		10.120.23.16
 ofc-prod-hms-c00.nix.sys.7d	A		10.120.23.20
 ofc-prod-hms-c01.nix.sys.7d	A		10.120.23.21
 ofc-prod-hms-d00.nix.sys.7d	A		10.120.23.18
 ofc-prod-hms-d01.nix.sys.7d	A		10.120.23.19
 prod-hms-http00.nix.sys.7d	A		10.112.6.44
 prod-hms-logstash00.nix.sys.7d	CNAME		prod-hms-monitoring00.nix.sys.7d
 prod-hms-metadata-storage00.nix.sys.7d	A		10.108.3.42
 prod-hms-metadata-storage01.nix.sys.7d	A		10.108.3.46
 prod-hms-storage00-meh.nix.sys.7d	A		10.120.23.11
 prod-hms-wr00.nix.sys.7d	A		10.120.23.12
 prod-hms-wr01-bonded.nix.sys.7d	A		10.120.23.15
 prod-hms-wr01.nix.sys.7d	A		10.120.23.13
 prod-hms-wr02.nix.sys.7d	A		10.120.23.22
 prod-hms00.nix.sys.7d	A		10.108.16.16
 prod-hms01.nix.sys.7d	A		10.112.16.16
 prod-hms02.nix.sys.7d	A		10.108.16.17
 prod-hms03.nix.sys.7d	A		10.108.16.21
 prod-hms04.nix.sys.7d	A		10.112.16.21

## Complex (cont. 2)

```
• root@ctr-prod-hms-b00:~# gluster vol info
Volume Name: gv01
Type: Distributed-Replicate
Volume ID: 94cff867-4a2c-4cc3-9a9b-8033623d49b6
Status: Started
Number of Bricks: 25 x 2 = 50
Transport-type: tcp
Bricks:
Brick1: ofc-prod-hms-b00:/export/brick000
Brick2: ctr-prod-hms-b00:/export/brick000
Brick3: ofc-prod-hms-b00:/export/brick001
Brick4: ctr-prod-hms-b00:/export/brick001
Brick5: ofc-prod-hms-b00:/export/brick002
Brick6: ctr-prod-hms-b00:/export/brick002
Brick7: ofc-prod-hms-b00:/export/brick003
Brick8: ctr-prod-hms-b00:/export/brick003
Brick9: ofc-prod-hms-b00:/export/brick004
Brick10: ctr-prod-hms-b00:/export/brick004
Brick11: gs2-prod-hms-a00:/export/brick000
Brick12: ofc-prod-hms-a00:/export/brick000
Brick13: gs2-prod-hms-a00:/export/brick001
Brick14: ofc-prod-hms-a00:/export/brick001
Brick15: gs2-prod-hms-a00:/export/brick002
```

## Complex (cont. 3)

```
Brick38: ofc-prod-hms-b01:/export/brick003
Brick39: ctr-prod-hms-b01:/export/brick004
Brick40: ofc-prod-hms-b01:/export/brick004
Brick41: gs2-prod-hms-a02:/export/brick000
Brick42: ofc-prod-hms-a02:/export/brick000
Brick43: gs2-prod-hms-a02:/export/brick001
Brick44: ofc-prod-hms-a02:/export/brick001
Brick45: gs2-prod-hms-a02:/export/brick002
Brick46: ofc-prod-hms-a02:/export/brick002
Brick47: gs2-prod-hms-a02:/export/brick003
Brick48: ofc-prod-hms-a02:/export/brick003
Brick49: gs2-prod-hms-a02:/export/brick004
Brick50: ofc-prod-hms-a02:/export/brick004
Options Reconfigured:
cluster.self-heal-daemon: off
cluster.data-self-heal: off
```

```
cluster.metadata-self-heal: off
cluster.entry-self-heal: off
```

## Complex (cont. 4)

```
• root@prod-hms01:~# netstat -natp| grep glusterfs | head
tcp        0      0 10.112.16.16:935      10.120.23.11:49182    ESTABLISHED 1917/gl
tcp        0      0 10.112.16.16:928      10.120.23.11:49186    ESTABLISHED 1917/gl
tcp        0      0 10.112.16.16:956      10.120.23.10:49155    ESTABLISHED 1917/gl
tcp        0      0 10.112.16.16:969      10.108.16.10:49185    ESTABLISHED 1917/gl
tcp        0      0 10.112.16.16:926      10.112.16.11:49153    ESTABLISHED 1917/gl
tcp        0      0 10.112.16.16:944      10.120.23.17:49155    ESTABLISHED 1917/gl
tcp        0      0 10.112.16.16:958      10.108.33.29:49154    ESTABLISHED 1917/gl
tcp        0      0 10.112.16.16:947      10.120.23.17:49152    ESTABLISHED 1917/gl
tcp        0      0 10.112.16.16:929      10.120.23.11:49185    ESTABLISHED 1917/gl
tcp        0      0 10.112.16.16:961      10.120.23.10:49152    ESTABLISHED 1917/gl
root@prod-hms01:~# netstat -natp| grep glusterfs | wc -l
51
```

## Issues

```
• ops@prod-hms03:~$ ls -l /srv/gv02/hms/track/000/049/009/850/17
-----T 1 root root 0 Nov  4  2015 /srv/gv02/hms/track/000/049/009/850/17
```

## Issues (cont.)

```
• Track: 45086687, 17
gv02: gs2-prod-hms-c00.nix.sys.7d
On 2 brick(s)
Result 0: -rw-rw-r-- 2 www-data www-data 5080970 May  5  2015 /export/brick002/gv02/hms
Brick brick002 owner: www-data size: 5080970
Result 1: -rw-rw-r-- 2 www-data www-data 5080970 May  5  2015 /export/brick003/gv02/hms
Brick brick003 owner: www-data size: 5080970
gv02: ctr-prod-hms-d00.nix.sys.7d
On 1 brick(s)
Result 0: -----T 2 root root 0 Jul 23 06:43 /export/brick001/gv02/hms/track/000/045
Brick brick001 owner: root size: 0
gv02: gs2-prod-hms-c01.nix.sys.7d
On 1 brick(s)
Result 0: -----T 2 root root 0 Dec 30  2015 /export/brick001/gv02/hms/track/000/045
Brick brick001 owner: root size: 0
```

```

gv02: ofc-prod-hms-d00.nix.sys.7d
On 1 brick(s)
Result 0: -----T 2 root root 0 Jul 23 06:43 /export/brick001/gv02/hms/track/000/045
Brick brick001 owner: root size: 0
gv02: ofc-prod-hms-c00.nix.sys.7d
On 2 brick(s)
Result 0: -rw-rw-r-- 2 www-data www-data 5080970 May 5 2015 /export/brick002/gv02/hms
Brick brick002 owner: www-data size: 5080970
Result 1: -rw-rw-r-- 2 www-data www-data 5080970 May 5 2015 /export/brick003/gv02/hms
Brick brick003 owner: www-data size: 5080970
gv02: ofc-prod-hms-c01.nix.sys.7d
On 1 brick(s)
Result 0: -----T 2 root root 0 Dec 30 2015 /export/brick001/gv02/hms/track/000/045
Brick brick001 owner: root size: 0
On 6 servers and 8 bricks, suggested actions:
ofc-prod-hms-c00.nix.sys.7d: cp /export/brick003/gv02/hms/track/000/045/086/687/17 /tmp
gs2-prod-hms-c00.nix.sys.7d: rm /export/brick002/gv02/hms/track/000/045/086/687/17
gs2-prod-hms-c00.nix.sys.7d: rm /export/brick003/gv02/hms/track/000/045/086/687/17
ctr-prod-hms-d00.nix.sys.7d: rm /export/brick001/gv02/hms/track/000/045/086/687/17
gs2-prod-hms-c01.nix.sys.7d: rm /export/brick001/gv02/hms/track/000/045/086/687/17
ofc-prod-hms-d00.nix.sys.7d: rm /export/brick001/gv02/hms/track/000/045/086/687/17
ofc-prod-hms-c00.nix.sys.7d: rm /export/brick002/gv02/hms/track/000/045/086/687/17
ofc-prod-hms-c00.nix.sys.7d: rm /export/brick003/gv02/hms/track/000/045/086/687/17
ofc-prod-hms-c01.nix.sys.7d: rm /export/brick001/gv02/hms/track/000/045/086/687/17
ofc-prod-hms-c01.nix.sys.7d: rm /srv/gv02/hms/track/000/045/086/687/17
ofc-prod-hms-c00.nix.sys.7d: install -g www-data -o www-data -m 664 /tmp/tmp.JD2j8fnkH
ofc-prod-hms-c00.nix.sys.7d: rm /tmp/tmp.JD2j8fnkHM

```

## Fixes

- 12:53:54-ccoffey@ciaranc:~/hms/code (master)\$ python ./gluster\_heal.py  
58 broken file(s) reported  
15 actually broken  
15 file(s) fixed, 0 not fixed

## Questions?